

Dongfang Zhao, Ph.D.

Nationality: United States of America

Education

University of Washington (UW), Seattle	Computer Science, PostDoc	2017
Illinois Institute of Technology, Chicago	Computer Science, PhD	2015
Emory University, Atlanta, GA	Computer Science, MS	2008
Katholieke Universiteit Leuven, Belgium	Statistics, Diploma	2005

Professional Experience

2026–now	Open-source Contributor, Meta/Facebook Faiss (CUDA kernel)
2025–now	Open-source Contributor, Microsoft DiskANN (C++ main branch)
2023–now	Assistant Professor, UW Tacoma School of Engineering & Technology
2023–now	Affiliate Faculty, UW eScience Institute
2018–2023	Visiting Professor, University of California, Davis
2018–2023	Assistant Professor, University of Nevada, Reno (UNR)
2015–2015	Research Intern, IBM Almaden Research Center, San Jose, CA
2014–2014	Research Intern, Argonne National Laboratory, Chicago, IL
2013–2013	Research Intern, Pacific Northwest National Laboratory, Seattle, WA
2009–2011	Software Developer (full-time), Epic Systems Corporation, Madison, WI

Selected Awards

2022	Federal Government Grant, UNR PI, NSF 2112345, \$255,916 “A DLT Machine Learning Platform for Blockchain Warehousing”
2020	Federal Government Grant, UNR PI, DOE SC0020455, \$200,000 “HPChain: Energy-Efficient Scalable Blockchains for HPC Systems”
2019	Best Paper Award, International Conference on Cloud Computing “BlockLite: Toward Accurate and Efficient Emulation of Public Blockchains in the Cloud”
2018	Best Student Paper Award, IEEE International Conference on Cloud Computing “Toward Cost-Effective Memory Scaling in Clouds: Symbiosis of Virtual and Physical Memory”
2016	Postdoctoral Fellowship, Sloan Foundation, \$155,000
2007	Graduate Fellowship, Oak Ridge Institute for Science and Education, \$85,000

Recent Projects

[EGA]	Euclidean geodesic alignment for self-limiting adapters of frozen encoders perserving out-of-distribution geometry in vector search: https://github.com/hpdc/EGA Tags: Representation learning, Parameter-efficient fine-tuning (PEFT), SigLIP, DINOv2
[ElasticIVF]	Elastic in-memory vector indexing and search for streaming data analysis on heterogeneous distributed platforms: https://github.com/hpdc/ElasticIVF Tags: Distributed GPU computing, Message passing interface (MPI), Faiss, CUDA kernel
[AdaDisk]	Decentralized agentic systems for adaptive ingestion-query scheduling of disk-resident vector indices in LLM serving: https://github.com/hpdc/AdaDisk Tags: Large language models, AI agents, Distributed systems, DiskANN, C++, Python

Selected Publications (Dongfang Zhao was/is the *student's academic advisor)

- [1] **Dongfang Zhao**. SIVF: GPU-Resident IVF Index for Streaming Vector Analytics. The 35th ACM International Symposium on High-Performance Parallel and Distributed Computing (HPDC), 2026.
- [2] *Jiuzhou Fu, **Dongfang Zhao**. QPAD: Quantile-Preserving Approximate Dimension Reduction for Nearest Neighbors Preservation in High-Dimensional Vector Search. The 42nd IEEE International Conference on Data Engineering (ICDE), 2026.
- [3] *Baigang Chen, **Dongfang Zhao**. Reliable Non-Leveled Homomorphic Encryption for Web Services. The ACM Web Conference (formerly International World Wide Web Conference, or WWW), 2026.
- [4] *Chengyu Gong, *Gefei Shen, Luanzheng Guo, Nathan Tallent, **Dongfang Zhao**. Order-Preserving Dimension Reduction for Multimodal Semantic Embedding. The 40th Annual AAAI Conference on Artificial Intelligence (AAAI), 2026.
- [5] *Olamide Timothy Tawose, Jun Dai, Lei Yang, **Dongfang Zhao**. Toward Efficient Homomorphic Encryption for Outsourced Databases through Parallel Caching. Proceedings of the ACM on Management of Data (SIGMOD), 2023.
- [6] *Abdullah Al-Mamun, Feng Yan, **Dongfang Zhao**. SciChain: Blockchain-enabled Lightweight and Efficient Data Provenance for Reproducible Scientific Computing. 37th IEEE International Conference on Data Engineering (ICDE), 2021.
- [7] *Abdullah Al-Mamun, Feng Yan, **Dongfang Zhao**. BAASH: lightweight, efficient, and reliable blockchain-as-a-service for HPC systems. The International Conference for High Performance Computing, Networking, Storage and Analysis (SC), 2021.
- [8] *Xinying Wang, *Olamide Timothy Tawose, Feng Yan, **Dongfang Zhao**. HDK: Toward High-Performance Deep-Learning-Based Kirchhoff Analysis. The Thirty-Fourth AAAI Conference on Artificial Intelligence (AAAI), 2020.
- [9] Heyang Qin, Syed Zawad, Yanqi Zhou, Lei Yang, **Dongfang Zhao**, Feng Yan. Swift machine learning model serving scheduling: a region based reinforcement learning approach. The International Conf. for High Performance Computing, Networking, Storage and Analysis (SC), 2019.
- [10] Parmita Mehta, Sven Dorkenwald, **Dongfang Zhao**, Tomer Kaftan, Alvin Cheung, Magdalena Balazinska, Ariel Rokem, Andrew J. Connolly, Jacob VanderPlas, Yusra AlSayyad. Comparative Evaluation of Big-Data Systems on Scientific Image Analytics Workloads. Proceedings of Very Large Data Bases (VLDB), 2017.
- [11] **Dongfang Zhao**, Kan Qiao, Jian Yin, Ioan Raicu. Dynamic Virtual Chunks: On Supporting Efficient Accesses to Compressed Scientific Data. IEEE Transactions on Services Computing (TSC). 9(1): 96-109 (2016).
- [12] **Dongfang Zhao**, Ning Liu, Dries Kimpe, Robert B. Ross, Xian-He Sun, Ioan Raicu. Towards Exploring Data-Intensive Scientific Applications at Extreme Scales through Systems and Simulations. IEEE Transactions on Parallel and Distributed Systems (TPDS). 27(6): 1824-1837 (2016).
- [13] **Dongfang Zhao**, Li Yang. Incremental Isometric Embedding of High-Dimensional Data Using Connected Neighborhood Graphs. IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI). 31(1): 86-98 (2009).